



CMM707 Cloud Computing

Coursework Report

MSc in Big Data Analytics

Submitted By

A.S.Fouzul Hassan

RGU No : 2410544 / IIT No : 20233214

Contents

1. System Architecture Diagram	3
<u>1.1.</u> Internal communication flow.....	4
2. Security & Ethical Challenges	5
3. CI/CD Process	6
4. Implementation	6
4.1. Microservices Implementation.....	6
4.2. Aggregator Service Implementation	10
4.3. CI/CD Process Implementation.....	13
4.4. CI/CD Deployment.....	15
5. Testing of the Deployed App	18
6. Observability.....	22
7. MoonInsurance System Deployment Runbook	24
7.1. Overview.....	24
7.2. Support Contacts	24
7.3. Process.....	25
8. Dashboard – MoonInsurance System	27
9. Metrics and High-Performance Analytics Support.....	28

1. System Architecture Diagram

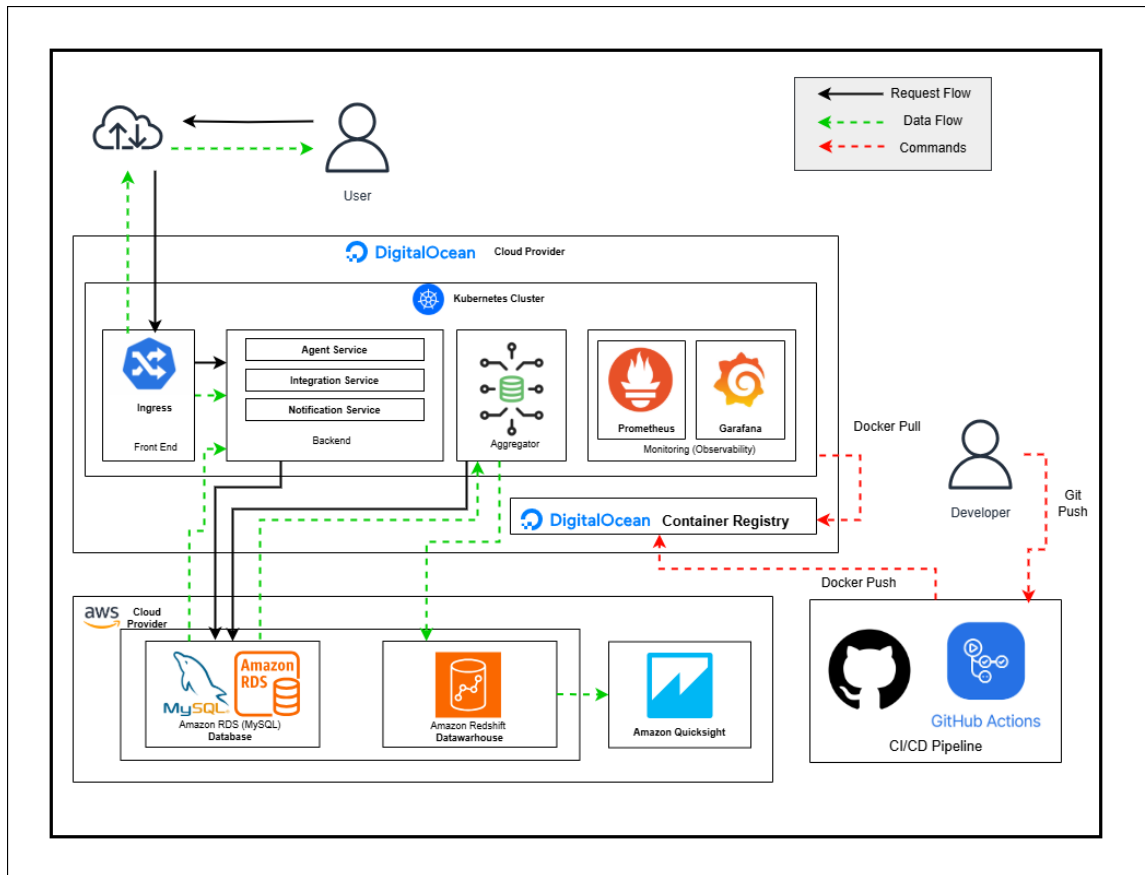


Figure 1- Solution Architecture Diagram for the System

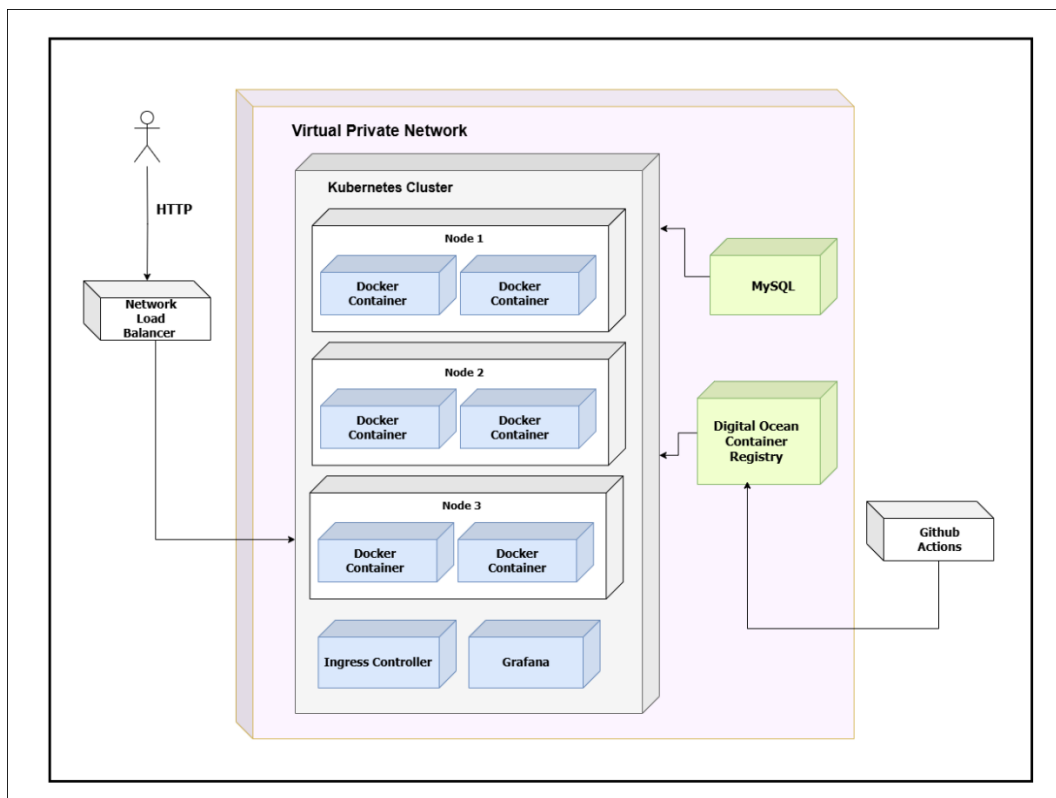


Figure 2 - Deployment Diagram

MoonInsurance is architected for scalability, security, fault-tolerance, and cost-efficiency. Docker images for every microservice are stored in the DigitalOcean Container Registry, while workloads are orchestrated by a managed DigitalOcean Kubernetes (DOKS) cluster.

- **Elastic scaling & HA** – Horizontal Pod Autoscalers, self-healing ReplicaSets and a regional Network Load Balancer keep services highly-available, while automatic node-pool scaling absorbs traffic spikes with zero-downtime Blue-Green releases.
- **Network & runtime isolation** – All service-to-service traffic is confined to a private VPC; an NGINX Ingress Controller terminates TLS and exposes only public APIs. Role-Based Access Control (RBAC), Kubernetes Secrets and DigitalOcean VPC firewalls gate access to runtime and data planes.
- **Observability** – Prometheus and Grafana (deployed via the Grafana Helm chart) stream metrics, logs and traces, raising alerts long before customers feel an impact.
- **Cost control** – DOKS pay-as-you-go pricing, spot-node pools for non-production workloads, and Open-Source monitoring tooling keep OpEx low.
- **Automated delivery** – A GitHub Actions pipeline builds, scans, and signs container images, pushes them to DOCR, then deploys them to a temporary *Green* namespace. Integration tests run in-cluster; on success traffic is flipped from *Blue* to *Green*. Rollback is instantaneous by re-pointing the Ingress.

1.1 Internal communication flow

1. External requests

User → Cloud Load Balancer → NGINX Ingress → microservice Pods (Agent, Integration, Notification).

2. Service discovery & east-west traffic

Pods talk to each other through ClusterIP services using Kubernetes DNS (agent-service.default.svc.cluster.local). mTLS can be toggled on with Linkerd if the compliance model tightens.

3. Data layer

- **MySQL RDS** on AWS stores transactional data for agents and sales.
- A scheduled **Aggregator CronJob** reads this data, computes KPIs (team sales, branch performance, product target attainment) and loads the results into **AWS Redshift Serverless**.

4. Analytics & dashboards

Redshift is the single source of truth for analytics; **AWS QuickSight** consumes the aggregated tables to surface real-time dashboards for executives.

5. CI/CD feedback loop

After every push, GitHub Actions → DOCR → DOKS. Deployment and test logs are shipped back to Grafana Loki, closing the observability loop.

This layered, modular communication pattern guarantees secure request routing, efficient data exchange, and hands-off operations—keeping MoonInsurance reliable as its user base and data volumes grow.

2. Security & Ethical Challenges

The **MoonInsurance** platform faces critical **security and ethical challenges** that must be addressed to maintain user trust, system integrity, and responsible data governance in the insurance domain. Since the platform processes sensitive information including policyholder identities, sales performance records, financial transactions, and team-level metrics, it becomes a potential target for cyber threats such as data breaches, credential theft, and ransomware.

From a **security** standpoint, the distributed microservices architecture—deployed via Kubernetes—introduces several vulnerabilities, especially during inter-service communication, API exposure, and data storage operations. Unauthorized access or data leakage during transmission could compromise customer confidentiality. To mitigate such risks, **end-to-end encryption** of data (both in-transit and at-rest), **API gateway rate limiting**, and **multi-factor authentication (MFA)** are enforced across all services. The use of **Role-Based Access Control (RBAC)** ensures agents, administrators, and developers can only access relevant information based on their roles.

Insider threats also pose a significant concern. Malicious or negligent access to sensitive sales or bonus data could disrupt fair agent compensation. This is mitigated through **auditing logs**, **user access reviews**, and **automated anomaly detection** integrated into the observability pipeline. All database connections are secured using environment-specific secrets stored securely through Kubernetes secrets management and GitHub Actions OIDC tokens for deployment.

On the **ethical** side, the MoonInsurance platform must handle **data ownership, fairness, and transparency** with care. Sales performance and aggregated metrics are used to make decisions about agent rankings, bonuses, and team performance—making it vital to ensure **data is not manipulated** or biased. The Aggregator microservice anonymizes data before analysis, and dashboards are permissioned in QuickSight to restrict visibility by role and region.

Furthermore, **algorithmic fairness** must be considered if predictive analytics or scoring systems are introduced. For instance, if the platform evolves to recommend training for underperforming agents based on AI predictions, it must ensure the model does not discriminate based on branch, gender, or historical bias. Similarly, insights derived from company-wide metrics should not be used for **commercial exploitation without consent** from agents or stakeholders.

The platform must remain transparent in how data is collected, processed, and used. Ethical safeguards such as **explicit consent flags**, **clear privacy notices**, and **data minimization** should be enforced. These practices must align with regulatory frameworks such as **Sri Lanka's Personal Data Protection Act (PDPA)** and other international data standards if the platform scales globally.

As a conclusion, MoonInsurance must adopt a **proactive and ethical-first approach** to both security and data governance. This includes:

- Implementing **robust security protocols**
- Maintaining **transparent ethical practices**
- Conducting **regular penetration tests and compliance audits**
- Enabling **real-time monitoring and alerts**
- Educating users on **privacy and responsible data handling**

Such measures not only improve resilience and reliability but also ensure the platform upholds high standards of **trust, fairness, and accountability** in the insurance technology sector.

3. CI/CD Process

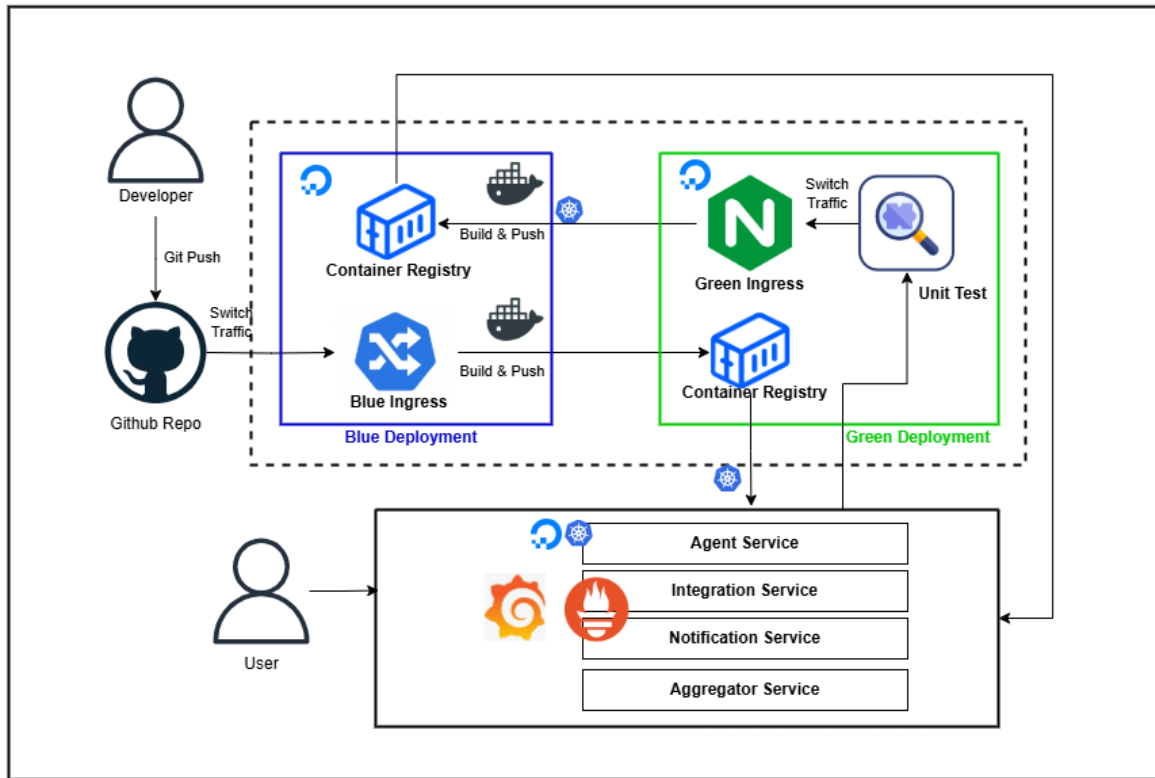


Figure 3 - Diagram for CI/CD Process

4. Implementation

4.1. Microservices Implementation

Initially a MySQL Database is been created in Amazon RDS and connected remotely to the MySQL Workbench.

Databases (1)							Group resources		Modify	Actions	Create database
Filter by databases											
DB identifier	Status	Role	Engine	Region ...	Size						
mooninsurance	Available	Instance	MySQL Co...	eu-north-1a	db.t4g.micro						

Figure 4 - Amazon RDS (MySQL)

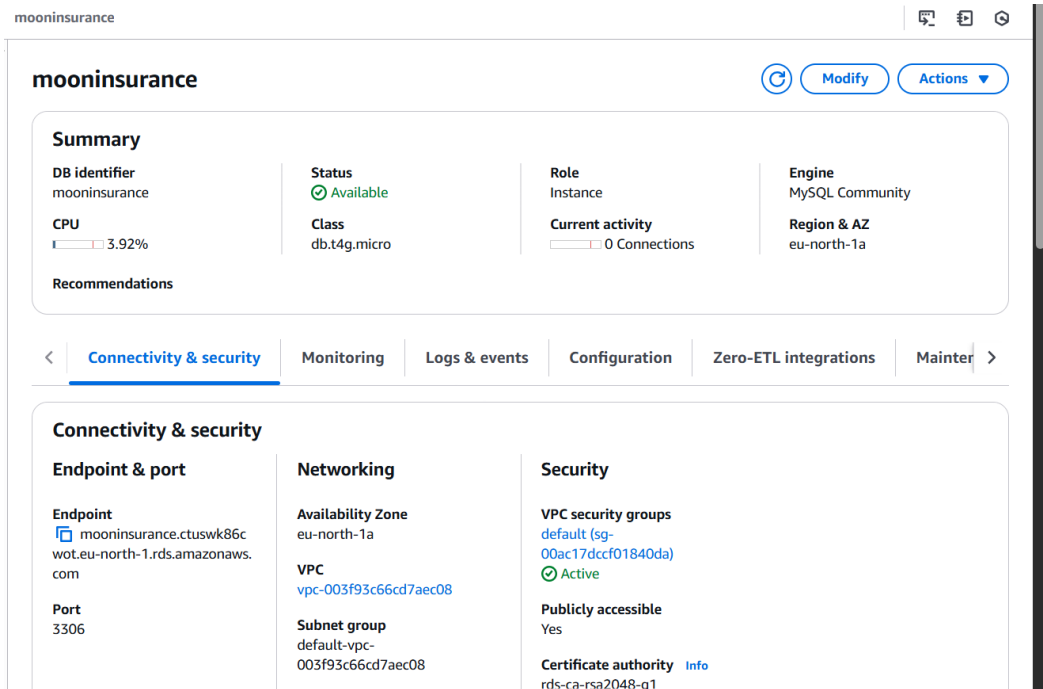


Figure 5 - Amazon RDS (MySQL) Details

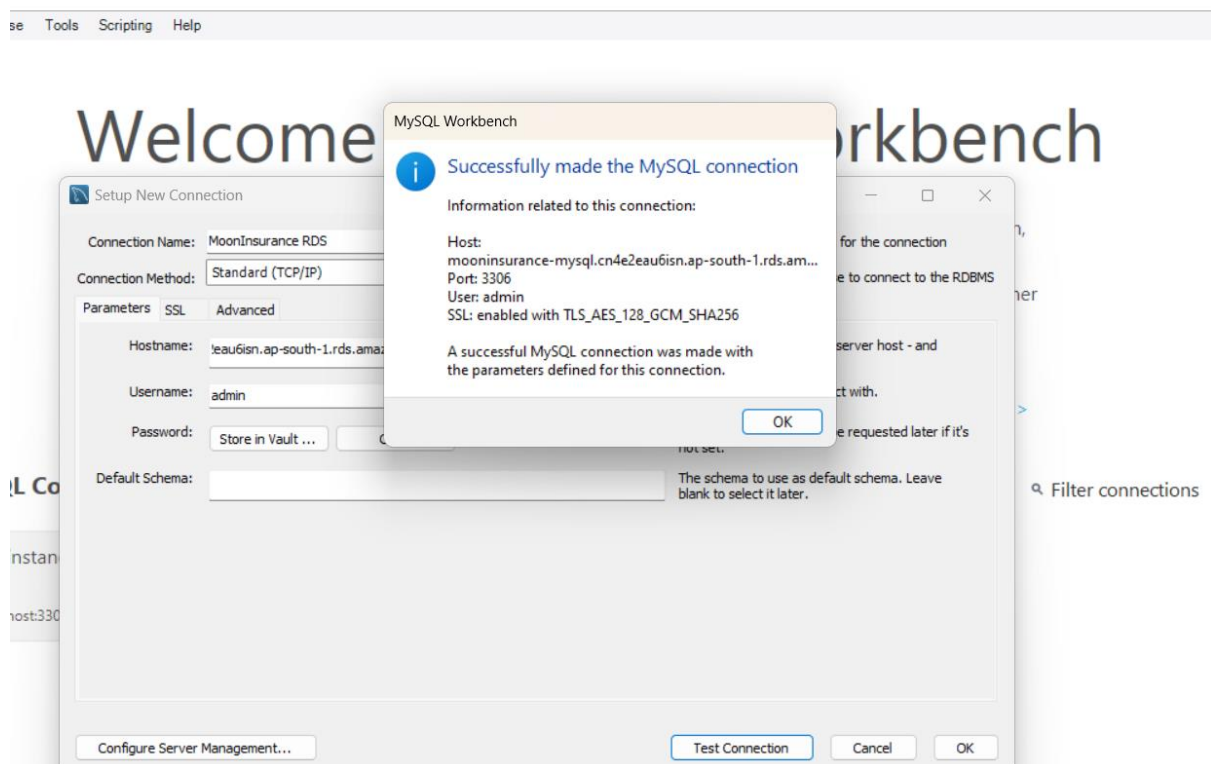


Figure 6 - Connecting AWS RDS to MySQL Workbench

Database is created and also few dummy data is inserted into the database.

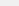
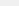
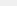
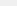
93 • <code>SELECT * FROM agents;</code>									
agent_id	agent_code	first_name	last_name	branch	contact_number	email			
10	AGT004	Anura	Jayasinghe	Colombo	0774567890	anura@example.com	20		
11	AGT005	Kamal	Gunawardena	Kandy	0775678901	kamal@example.com	20		
12	AGT006	Sunil	Dias	Galle	0776789012	sunil@example.com	20		


94 • <code>SELECT * FROM products;</code>									
product_id	product_type	product_name	target_amount						
1	Life	Life Insurance - Platinum Plan	120000.00						
2	Vehicle	Vehicle Insurance - Full Cover	90000.00						
3	Health	Health Insurance - Family Plan	100000.00						

95 • `SELECT * FROM sales;`

Result Grid

Filter Rows:

Edit:    





Export/Import: 


	sale_id	agent_code	amount	sale_date	branch	product_id
▶	5	AGT004	125000.00	2024-03-01	Colombo	1
	6	AGT005	92000.00	2024-03-02	Kandy	2
	7	AGT006	100000.00	2024-03-05	Galle	3

68 • `SELECT * FROM notifications;`

Result Grid

Filter Rows:

Edit:    

Export/Import: 

	id	agent_code	message	sent_at
▶	1	AGT001	* Nimal has reached 150k sales milestone!	2025-04-13 08:16:36
	2	AGT002	* Reminder: Vehicle sales target due this week.	2025-04-13 08:16:36
	3	AGT003	* Ruwan's monthly health insurance target ac...	2025-04-13 08:16:36
	4	AGT001	Test notification message	2025-04-13 08:18:36

Figure 7 - Few Screenshots of Implemented Tables

Microservice code API is developed with Python Flask.

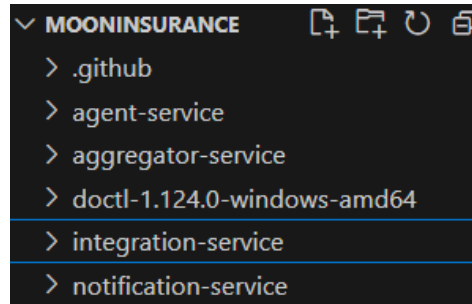


Figure 8 - File Structure of the solution

```

agent-service > agent_app.py > ...
1  from flask import Flask, request, jsonify, render_template
2  import mysql.connector
3
4  app = Flask(__name__)
5
12
13 # Route for the front-end page
14 @app.route('/agent')
15 def index():
16     return render_template('index.html')
17
18 # API to add a new agent
19 @app.route('/agent/add', methods=['POST'])
20 def add_agent():
21     data = request.json
22     try:
23         connection = mysql.connector.connect(
24             host=MYSQL_HOST,
25             user=MYSQL_USER,
26             password=MYSQL_PWD,
27             database=MYSQL_DB,
28             port=PORT
29         )
30         cursor = connection.cursor()
31
32         query = """

```

Figure 9 - Agent Services-Code Implementation


```

integration-service > integration_app.py > ...
12
13 @app.route('/integration')
14 def index():
15     return render_template('index.html')
16
17 # CREATE: Add a new sale
18 @app.route('/integration/add', methods=['POST'])
19 def add_sale():
20     data = request.json
21     print("Received Agent Data:", data)
22     try:
23         connection = mysql.connector.connect(
24             host=MYSQL_HOST, user=MYSQL_USER,
25             password=MYSQL_PWD, database=MYSQL_DB, port=PORT
26         )
27         cursor = connection.cursor()
28
29         query = """
30         INSERT INTO sales (agent_code, product_id, amount, sale_date, branch)
31         VALUES (%s, %s, %s, %s, %s)
32         """
33         cursor.execute(query, (
34             data.get('agent_code'),
35             data.get('product_id'),
36             data.get('amount'),
37             data.get('sale_date'),
38             data.get('branch')
39         ))
40         connection.commit()
41         return jsonify({"message": "Sale record added successfully."}), 201
42
43 except mysql.connector.Error as err:

```

Figure 10 - Integration Service Code-Implementation

```

notification-service > notification_app.py > send_notification
13 @app.route('/notification')
14 def index():
15     return render_template('index.html')
16
17 # Send a notification
18 @app.route('/notification/send', methods=['POST'])
19 def send_notification():
20     data = request.json
21     agent_code = data.get("agent_code")
22     message = data.get("message")
23
24     try:
25         connection = mysql.connector.connect(
26             host=MYSQL_HOST,
27             user=MYSQL_USER,
28             password=MYSQL_PWD,
29             database=MYSQL_DB,
30             port=PORT
31         )
32         cursor = connection.cursor()
33         query = "INSERT INTO notifications (agent_code, message) VALUES (%s, %s)"
34         cursor.execute(query, (agent_code, message))
35         connection.commit()
36         return jsonify({"message": f"Notification sent to agent {agent_code}."}), 201
37
38     except mysql.connector.Error as err:
39         print(f"Error: {err}")
40         return jsonify({"message": "Failed to send notification."}), 500
41
42 # Get all notifications
43 @app.route('/notification/get', methods=['GET'])
44 def get_notifications():

```

Figure 11 - Notification Service Implmentation

4.2. Aggregator Service Implementation

AWS Redshift Namespace and workgroups are created.

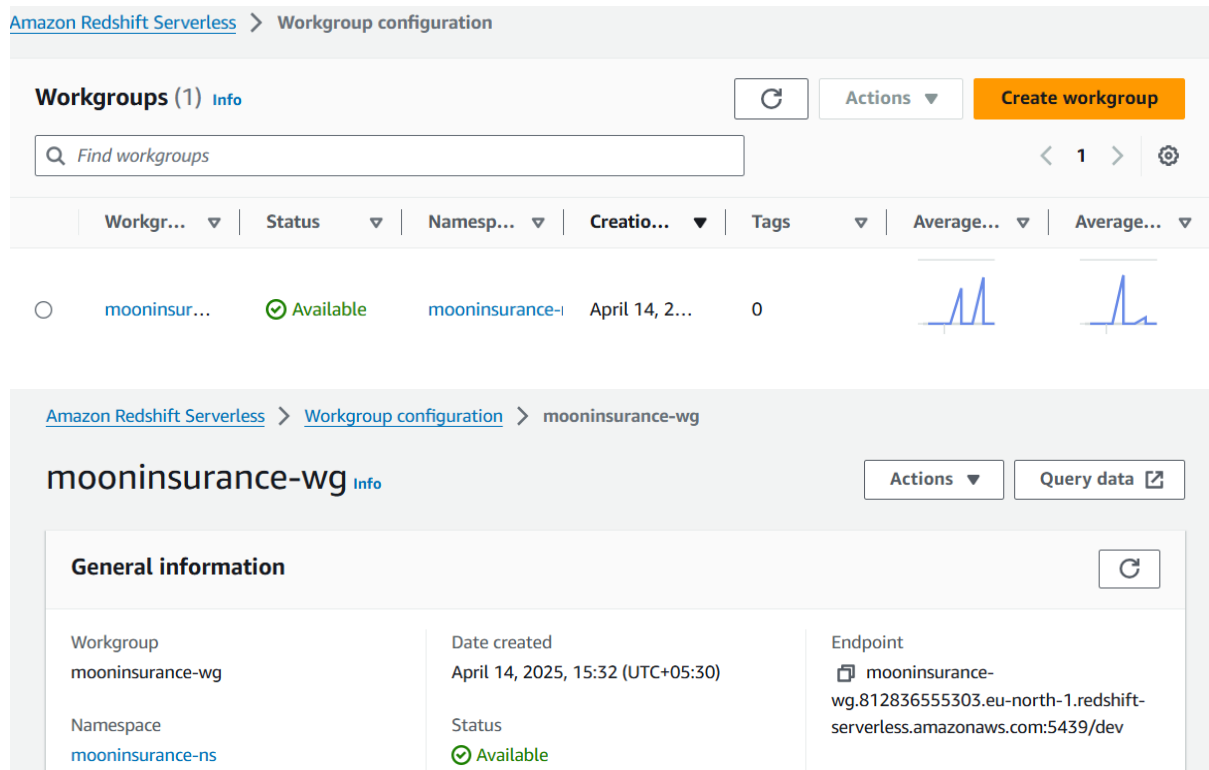
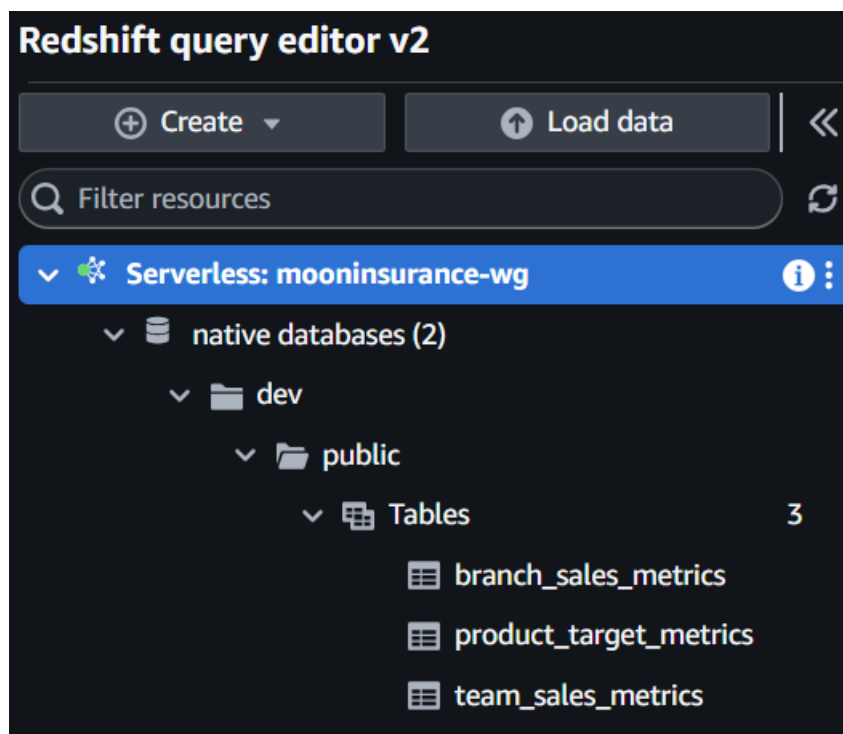


Figure 12 - AWS-Redshift

A Data Warehouse is created in Redshift for Aggregation Service.



An API is developed for aggregation, where the codes writes the data to the Datawarehouse.

Redshift query editor v2

Filter resources

Serverless: mooninsurance-wg

native databases (2)

dev

public

Tables

branch_sales_metrics

product_target_metrics

team sales metrics

branch_sales_metrics

Field	Type	NL	CMP
A branch	character varying(100)	NN	lzo
# total_sales	numeric(12,2)	NULL	az64

Run

Limit 100

Explain

Isolated session

Serverless: m... dev

Schedule

```

1 SELECT *
2 FROM
3 "dev"."public"."branch_sales_metrics";
4

```

Result 1 (6)

branch	total_sales
Colombo	212000
Kandy	191000
Anuradhapura	130000
Galle	100000
Matara	61000
Jaffna	32000

Redshift query editor v2

Filter resources

Serverless: mooninsurance-wg

native databases (2)

dev

public

Tables

branch_sales_metrics

product_target_metrics

team sales metrics

product_target_metrics

Field	Type	NL	CMP
A product_name	character varying(100)	NN	lzo
# target_amount	numeric(12,2)	NULL	az64
# total_sales	numeric(12,2)	NULL	az64
A status	character varying(20)	NULL	lzo

Run

Limit 100

Explain

Isolated session

Serverless: m... dev

Schedule

```

1 SELECT * FROM "dev"."public"."product_target_metrics";

```

Result 1 (5)

product_name	target_amount	total_sales	status
Life Insurance - Platinum ...	120000	255000	Achieved
Health Insurance - Family...	100000	199000	Achieved
Vehicle Insurance - Full C...	90000	179000	Achieved
Life Insurance - Basic Plan	60000	61000	Achieved
Vehicle Insurance - Third ...	30000	32000	Achieved

Redshift query editor v2

Filter resources

Serverless: mooninsurance-wg

native databases (2)

dev

public

Tables

branch_sales_metrics

product_target_metrics

team sales metrics

team_sales_metrics

Field	Type	NL	CMP
A team	character varying(50)	NN	lzo
# total_sales	numeric(12,2)	NULL	az64

Run

Limit 100

Explain

Isolated session

Serverless: m... dev

Schedule

```

1 SELECT * FROM "dev"."public"."team_sales_metrics";

```

Result 1 (4)

team	total_sales
Beta	222000
Gamma	187000
Delta	160000
Alpha	157000

Figure 13 - Redshift Query Editor (Datawarehouse)

```

def fetch_dataframe(query):
    try:
        conn = mysql.connector.connect(**MYSQL_CONFIG)
        df = pd.read_sql(query, conn)
        conn.close()
        return df
    except mysql.connector.Error as err:
        print(f"[MySQL ERROR] {err}")
        return pd.DataFrame()

def get_team_sales():
    return fetch_dataframe("""
        SELECT a.team, SUM(s.amount) AS total_sales
        FROM agents a
        JOIN sales s ON a.agent_code = s.agent_code
        GROUP BY a.team
        ORDER BY total_sales DESC
    """)

def get_product_target_achievement():
    return fetch_dataframe("""
        SELECT p.product_name, p.target_amount,
        COALESCE(SUM(s.amount), 0) AS total_sales,
        CASE
            WHEN COALESCE(SUM(s.amount), 0) >= p.target_amount
            THEN 'Achieved' ELSE 'Not Achieved'
        END AS status
        FROM products p
        LEFT JOIN sales s ON p.product_id = s.product_id
        GROUP BY p.product_id
    """)

```

Figure 14 - Fetch Aggregations

```

def connect_redshift():
    return psycopg2.connect(**REDSHIFT_CONFIG)

def upsert_team_sales(df):
    conn = connect_redshift()
    cursor = conn.cursor()
    for _, row in df.iterrows():
        cursor.execute("""
            DELETE FROM team_sales_metrics WHERE team = %s;
            """, (row['team'],))
        cursor.execute("""
            INSERT INTO team_sales_metrics (team, total_sales)
            VALUES (%s, %s);
            """, (row['team'], row['total_sales']))
    conn.commit()
    cursor.close()
    conn.close()
    print("✅ Upserted team_sales_metrics")

def upsert_product_target_metrics(df):
    conn = connect_redshift()
    cursor = conn.cursor()
    for _, row in df.iterrows():
        cursor.execute("""
            DELETE FROM product_target_metrics WHERE product_name = %s;
            """, (row['product_name'],))
        cursor.execute("""
            INSERT INTO product_target_metrics (product_name, target_amount, total_sales,
            VALUES (%s, %s, %s, %s);
            """, (row['product_name'], row['target_amount'], row['total_sales'], row['status']))
    conn.commit()

def upsert_branch_sales(df):
    conn = connect_redshift()
    cursor = conn.cursor()
    for _, row in df.iterrows():
        cursor.execute("""
            DELETE FROM branch_sales_metrics WHERE branch = %s;
            """, (row['branch'],))
        cursor.execute("""
            INSERT INTO branch_sales_metrics (branch, total_sales)
            VALUES (%s, %s);
            """, (row['branch'], row['total_sales']))
    conn.commit()
    cursor.close()
    conn.close()
    print("✅ Upserted branch_sales_metrics")

```

Figure 15 - Insert into Redshift

4.3. CI/CD Process Implementation

```
# Set working directory
WORKDIR /app

# Copy the service files
COPY . .

RUN pip install --no-cache-dir -r requirements.txt

EXPOSE 5000

CMD ["gunicorn", "-b", "0.0.0.0:5000", "agent_app:app"]
```

Figure 16 - Docker File

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: agent-service-blue
5  spec:
6    replicas: 1
7    selector:
8      matchLabels:
9        app: agent-service
10       version: blue
11    template:
12      metadata:
13        labels:
14          app: agent-service
15          version: blue
16      spec:
17        containers:
18          - name: agent-service
19            image: registry.digitalocean.com/moonregistry/agent-service:blue-v2
20            ports:
21              - containerPort: 80
22            imagePullSecrets:
23              - name: do-secret
24
25 ---
26
27 apiVersion: v1
28 kind: Service
29 metadata:
30   name: agent-service-blue
31 spec:
```

Figure 17 - Blue Ingress Deployment (Yaml)

```
io.k8s.api.core.v1.Service (v1@service.json) | io.k8s.api.apps.v1.Deployment (v1@deployment.json)
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: agent-service-green
5  spec:
6    replicas: 1
7    selector:
8      matchLabels:
9        app: agent-service
10       version: green
11    template:
12      metadata:
13        labels:
14          app: agent-service
15          version: green
16      spec:
17        containers:
18          - name: agent-service
19            image: registry.digitalocean.com/moonregistry/agent-service:green-v2
20            ports:
21              - containerPort: 80
22            imagePullSecrets:
23              - name: do-secret
24
25 ---
26
27 apiVersion: v1
28 kind: Service
29 metadata:
30   name: agent-service-green
31 spec:
```

Figure 18 - Green Ingress Deployment (Yaml)

Similarly, this is implemented for all the three services.

```

io.k8s.api.batch.v1.CronJob (v1@cronjob.json)
1  apiVersion: batch/v1
2  kind: CronJob
3  metadata:
4    name: aggregator-cronjob
5  spec:
6    schedule: "30 18 * * *" # Runs every day at 6:30 PM UTC (12 AM in Sri Lanka)
7    jobTemplate:
8      spec:
9        template:
10         spec:
11           containers:
12             - name: aggregator
13               image: registry.digitalocean.com/meditrackcontainer/aggregator-service:latest
14               imagePullPolicy: Always
15               restartPolicy: Never
16               imagePullSecrets:
17                 - name: do-secret
18

```

Figure 19 - Cron-job of AggregationService

```

! ingress-blue.yaml > {} spec > {} rules > {} 0 > {} http > {} paths > {}
io.k8s.api.networking.v1.Ingress (v1@ingress.json)
1  apiVersion: networking.k8s.io/v1
2  kind: Ingress
3  metadata:
4    name: microservices-ingress
5    annotations:
6      nginx.ingress.kubernetes.io/rewrite-target:
7  spec:
8    ingressClassName: nginx
9    rules:
10     - host: 209.38.124.165.nip.io
11       http:
12         paths:
13           - path: /agent
14             pathType: Exact
15             backend:
16               service:
17                 name: agent-service-blue
18                 port:
19                   number: 80
20           - path: /agent/get
21             pathType: Exact
22             backend:
23               service:
24                 name: agent-service-blue
25                 port:
26                   number: 80
27           - path: /agent/add
28             pathType: Exact
29             backend:
30               service:
31                 name: agent-service-blue

```

```

! ingress-green.yaml > {} spec > {} rules > {} 0 > {} http > {} paths > {}
io.k8s.api.networking.v1.Ingress (v1@ingress.json)
1  apiVersion: networking.k8s.io/v1
2  kind: Ingress
3  metadata:
4    name: microservices-ingress
5    annotations:
6      nginx.ingress.kubernetes.io/rewrite-target:
7  spec:
8    ingressClassName: nginx
9    rules:
10     - host: 209.38.124.165.nip.io
11       http:
12         paths:
13           - path: /agent
14             pathType: Exact
15             backend:
16               service:
17                 name: agent-service-green
18                 port:
19                   number: 80
20           - path: /agent/get
21             pathType: Exact
22             backend:
23               service:
24                 name: agent-service-green
25                 port:
26                   number: 80
27           - path: /agent/add
28             pathType: Exact
29             backend:
30               service:
31                 name: agent-service-green

```

Figure 20 - Green and Blue Ingress-YAML

4.4. CI/CD Deployment

CI (Build & Test):

Upon each code push to GitHub, the system automatically triggers GitHub Actions to build Docker images for all microservices (Agent, Integration, Notification, Aggregator) and execute automated tests to validate them.

```
on:
  push:
    branches:
      - main
```

CD (Deploy):

After a successful build and test phase, updated services are deployed to the DigitalOcean Kubernetes cluster.

- Deployment begins with the **Green** environment.
- Stability tests are conducted post-deployment.
- Traffic is switched from **Blue** to **Green** following the **Blue-Green deployment strategy**.
- The kubectl rollout restart command is used to restart the deployments and apply updates.
- The **Aggregator Service** is deployed as a **CronJob**, enabling periodic data aggregation and transfer to Redshift.

Actual Workflow (from YAML file)

1. change-to-blue : Build & push aggregator & Apply cronjob and ingress-blue

```
jobs:
  change-to-blue:
    runs-on: ubuntu-latest

    steps:
      - name: Checkout code
        uses: actions/checkout@v3

      - name: Log in to DigitalOcean
        run: echo "${{ secrets.DIGITALOCEAN_ACCESS_TOKEN }}" | docker login registry.digitalocean.com

      - name: Build and push Aggregator Service
        run: |
          docker build --no-cache -t registry.digitalocean.com/moonregistry/aggregator-service
          docker push registry.digitalocean.com/moonregistry/aggregator-service

      - name: Set up doctl
        uses: digitalocean/action-doctl@v2
        with:
          token: "${{ secrets.DIGITALOCEAN_ACCESS_TOKEN }}"

      - name: Install kubectl
        uses: azure/setup-kubectl@v3
```

2. build-green: Build & push agent, integration, notification (green)

```
✓ deploy-green:
  needs: build-green
  runs-on: ubuntu-latest

  steps:
    ✓ - name: Checkout code
      | uses: actions/checkout@v3

    ✓ - name: Set up doctl
      | uses: digitalocean/action-doctl@v2
      | with:
      |   token: ${ secrets.DIGITALOCEAN_ACCESS_TOKEN }

    ✓ - name: Install kubectl
      | uses: azure/setup-kubectl@v3
      | with:
      |   version: 'latest'

    ✓ - name: Configure kubeconfig
      | run: doctl kubernetes cluster kubeconfig save mooninsurance-cluster

    ✓ - name: Apply Green Deployments
      | run: |
      |   kubectl apply -f ./agent-service/agent-green-deployment.yaml
      |   kubectl apply -f ./integration-service/integration-green-deployment.yaml
      |   kubectl apply -f ./notification-service/notification-green-deployment.yaml
```

3. deploy-green: Deploy all green services to K8s using kubectl

```
test-green:
  needs: deploy-green
  runs-on: ubuntu-latest

  steps:
    - name: Wait for green deployments to stabilize
      | run: sleep 30

    - name: Test Agent Service (Green)
      | run: |
      |   curl -f http://209.38.124.165.nip.io/agent/get || exit 1

    - name: Test Integration Service (Green)
      | run: |
      |   curl -f http://209.38.124.165.nip.io/integration/get || exit 1

    - name: Test Notification Service (Green)
      | run: |
      |   curl -f http://209.38.124.165.nip.io/notification/get || exit 1
```

4. build-blue: Prepare next version of blue (agent/integration/notification)

```
✓ deploy-blue:
  needs: build-blue
  runs-on: ubuntu-latest

  steps:
    ✓ - name: Checkout code
      | uses: actions/checkout@v3

    ✓ - name: Set up doctl
      | uses: digitalocean/action-doctl@v2
      | with:
      |   token: ${ secrets.DIGITALOCEAN_ACCESS_TOKEN }

    ✓ - name: Install kubectl
      | uses: azure/setup-kubectl@v3
      | with:
      |   version: 'latest'

    ✓ - name: Configure kubeconfig
      | run: doctl kubernetes cluster kubeconfig save mooninsurance-cluster

    ✓ - name: Apply Blue Deployments and Ingress
      | run: |
      |   kubectl apply -f ./agent-service/agent-blue-deployment.yaml
      |   kubectl apply -f ./integration-service/integration-blue-deployment.yaml
      |   kubectl apply -f ./notification-service/notification-blue-deployment.yaml
      |   kubectl apply -f ingress-green.yaml
      |   kubectl rollout restart deployment agent-service-blue
      |   kubectl rollout restart deployment integration-service-blue
      |   kubectl rollout restart deployment notification-service-blue
```


5. deploy-blue: Restart blue deployments & Switch ingress if needed

```
deploy-blue:
  needs: build-blue
  runs-on: ubuntu-latest

  steps:
    - name: Checkout code
      uses: actions/checkout@v3

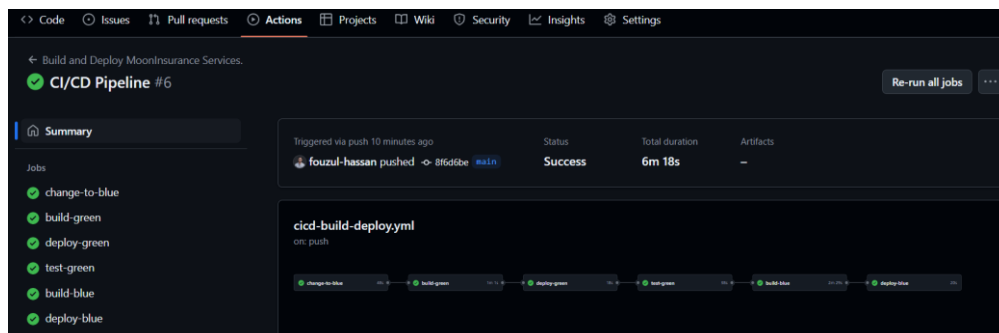
    - name: Set up doctl
      uses: digitalocean/action-doctl@v2
      with:
        token: ${ secrets.DIGITALOCEAN_ACCESS_TOKEN }

    - name: Install kubectl
      uses: azure/setup-kubectl@v3
      with:
        version: 'latest'

    - name: Configure kubeconfig
      run: doctl kubernetes cluster kubeconfig save mooninsurance-cluster
```

Blue-Green Deployment : These are two versions (Blue & Green) run in parallel which only one handles **live traffic**. The **inactive ones are been deployed**, test it, then switch

This ensures there is no any downtime, enables Automated builds and deployments, Safe rollback and Observability.



	agent-service 7 Images	Recent image tag: green	34 seconds ago	...
	aggregator-service 4 Images	Recent image tag: latest	1 minute ago	...
	integration-service 6 Images	Recent image tag: green	21 seconds ago	...
	notification-service 5 Images	Recent image tag: green	8 seconds ago	...

```
PS C:\Msc Files\Cloud Computing\CW\Solution\MoonInsurance> kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
agent-service-blue-5d848ffffb-ns65x 1/1     Running   0           5h34m
agent-service-green-56b47787cf-rw7q8 1/1     Running   0           5h36m
integration-service-blue-856ffc86c-sgz55 1/1     Running   0           5h34m
integration-service-green-7cd466d8c7-kh8qw 1/1     Running   0           5h36m
notification-service-blue-846699774f-cdcpj 1/1     Running   0           5h34m
notification-service-green-84cb8ffd6f-s9nhb 1/1     Running   0           5h36m
```

5. Testing of the Deployed App

Testing the deployed app using Postman.

1. Agent Service

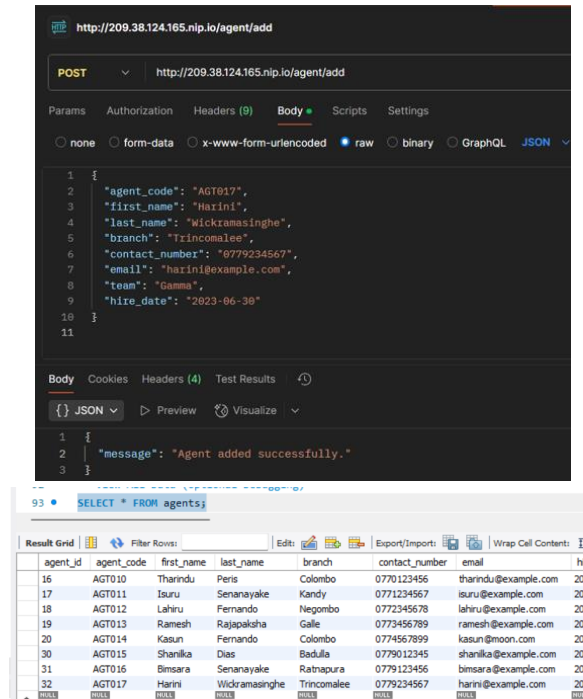


Figure 21 - Insert New-Agent and Verify

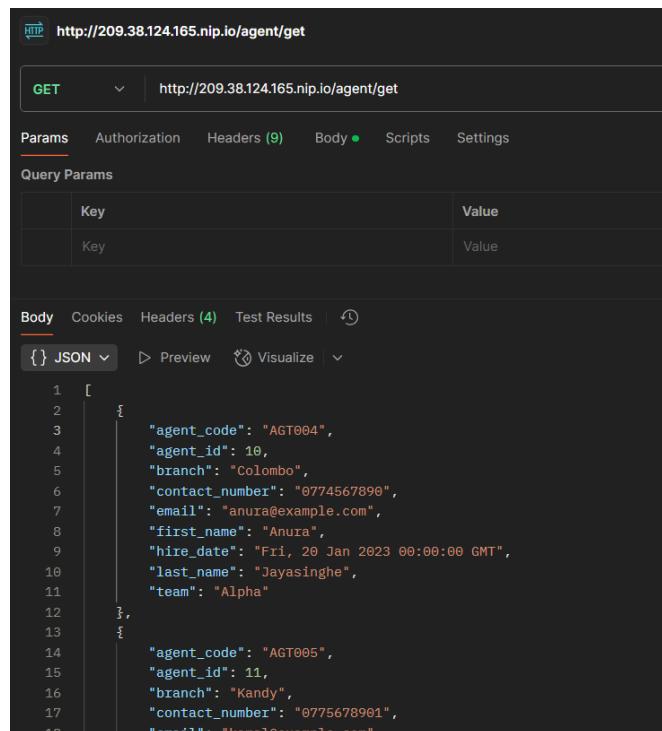


Figure 22-View-All Agents

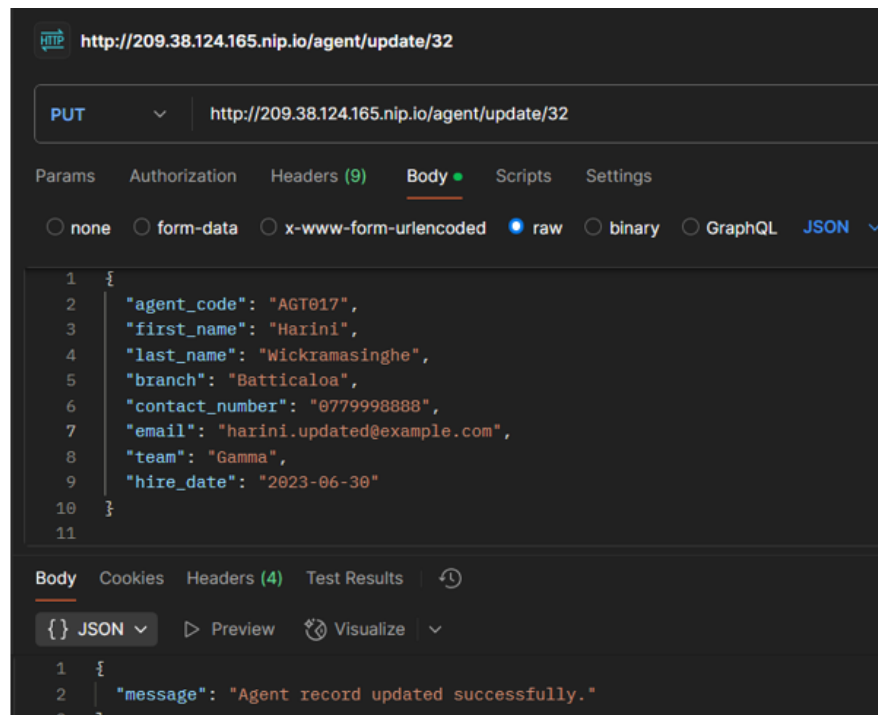


Figure 23 - Update-Agent

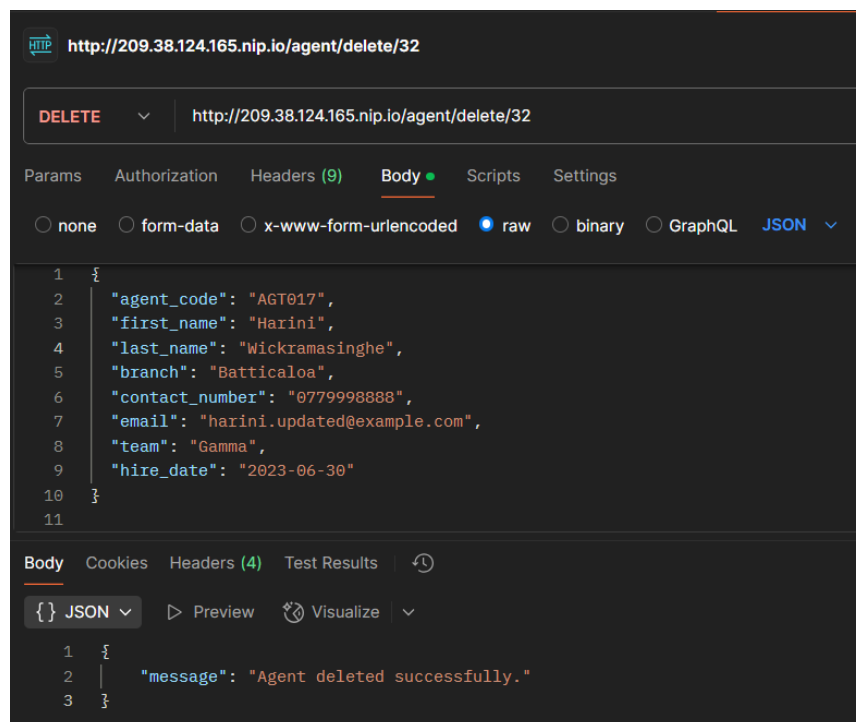


Figure 24 - Delete Agent

2. Integration Service

The first screenshot shows a REST client interface with a POST request to `http://209.38.124.165.nip.io/integration/add`. The request body is a JSON object:

```
{  "agent_code": "AGT013",  "product_id": 1,  "amount": 100000,  "sale_date": "2024-04-10",  "branch": "Colombo"}
```

The response body is a JSON object:

```
{  "message": "Sales record added successfully."}
```

The second screenshot shows a SQL query `SELECT * FROM sales;` with the following result grid:

	sale_id	agent_code	amount	sale_date	branch	product_id
▶	5	AGT004	125000.00	2024-03-01	Colombo	1
	6	AGT005	92000.00	2024-03-02	Kandy	2
	7	AGT006	100000.00	2024-03-05	Galle	3
	8	AGT007	61000.00	2024-03-06	Matara	4
	9	AGT008	32000.00	2024-03-07	Jaffna	5
	10	AGT009	130000.00	2024-03-08	Anuradhapura	1
	17	AGT015	58000.00	2024-03-09	Badulla	4
	18	AGT016	89000.00	2024-03-10	Ratnapura	2
	19	AGT017	104000.00	2024-03-11	Trincomalee	3

Figure 25 - Add-New Integration/Sales & Verify

The screenshot shows a REST client interface with a GET request to `http://209.38.124.165.nip.io/integration/get`. The response body is a JSON array of sales records:

```
[  {    "agent_code": "AGT004",    "amount": "125000.00",    "branch": "Colombo",    "product_id": 1,    "sale_date": "Fri, 01 Mar 2024 00:00:00 GMT",    "sale_id": 5  },  {    "agent_code": "AGT005",    "amount": "92000.00",    "branch": "Kandy",    "product_id": 2,    "sale_date": "Sat, 02 Mar 2024 00:00:00 GMT",    "sale_id": 6  }]
```

Figure 26 - View-all Integration/Sales

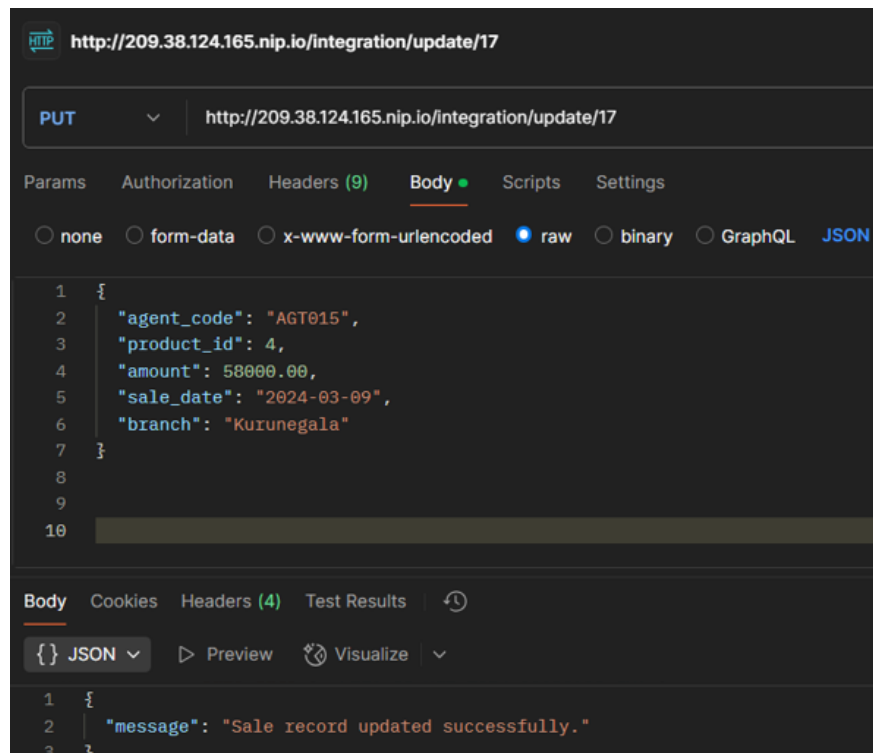


Figure 27 - Update-Integration/Sales

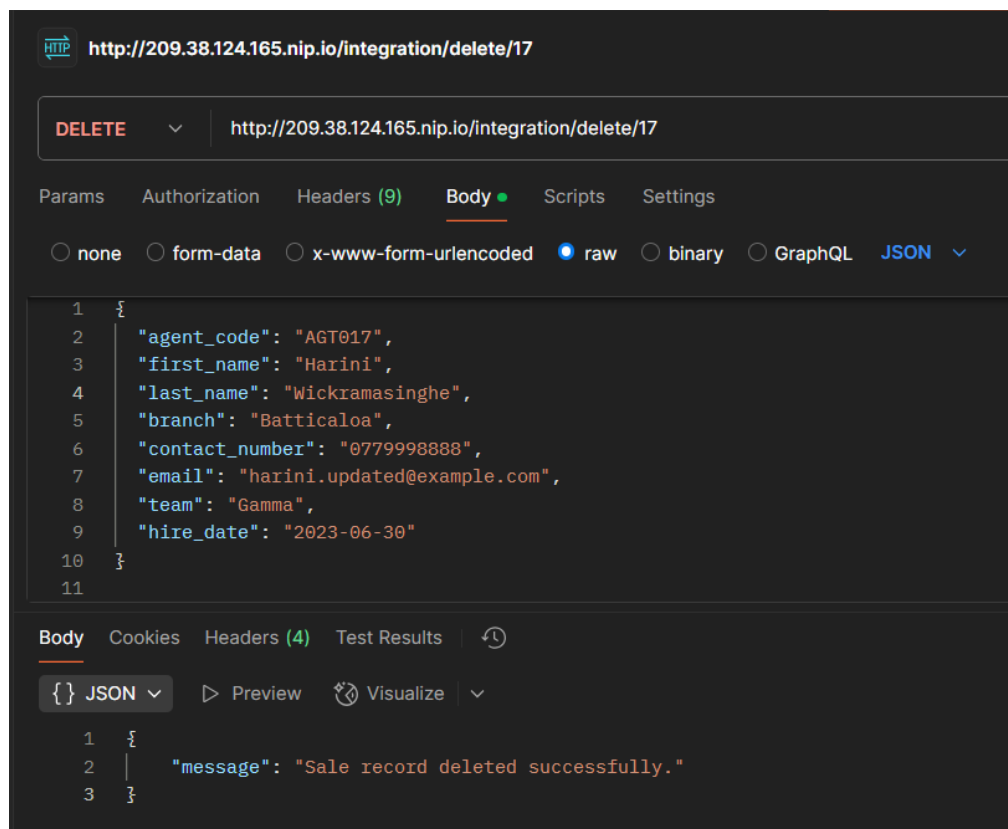
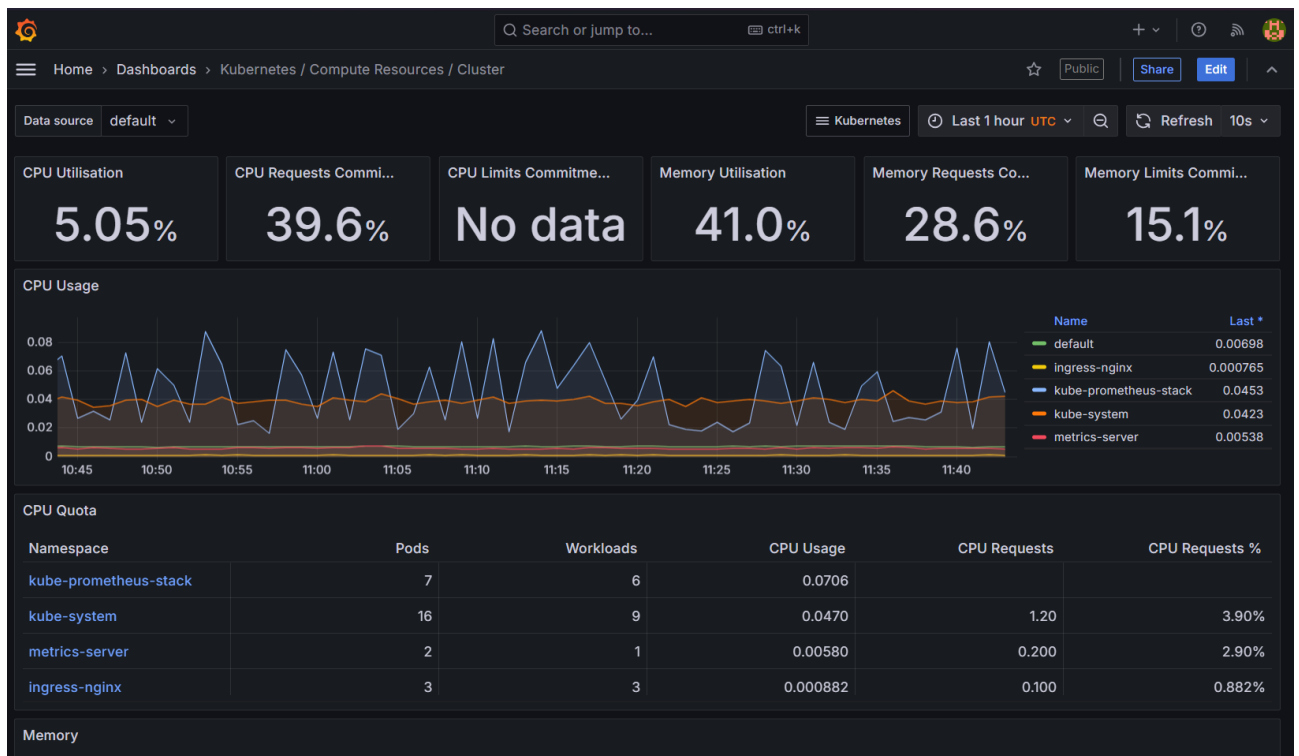
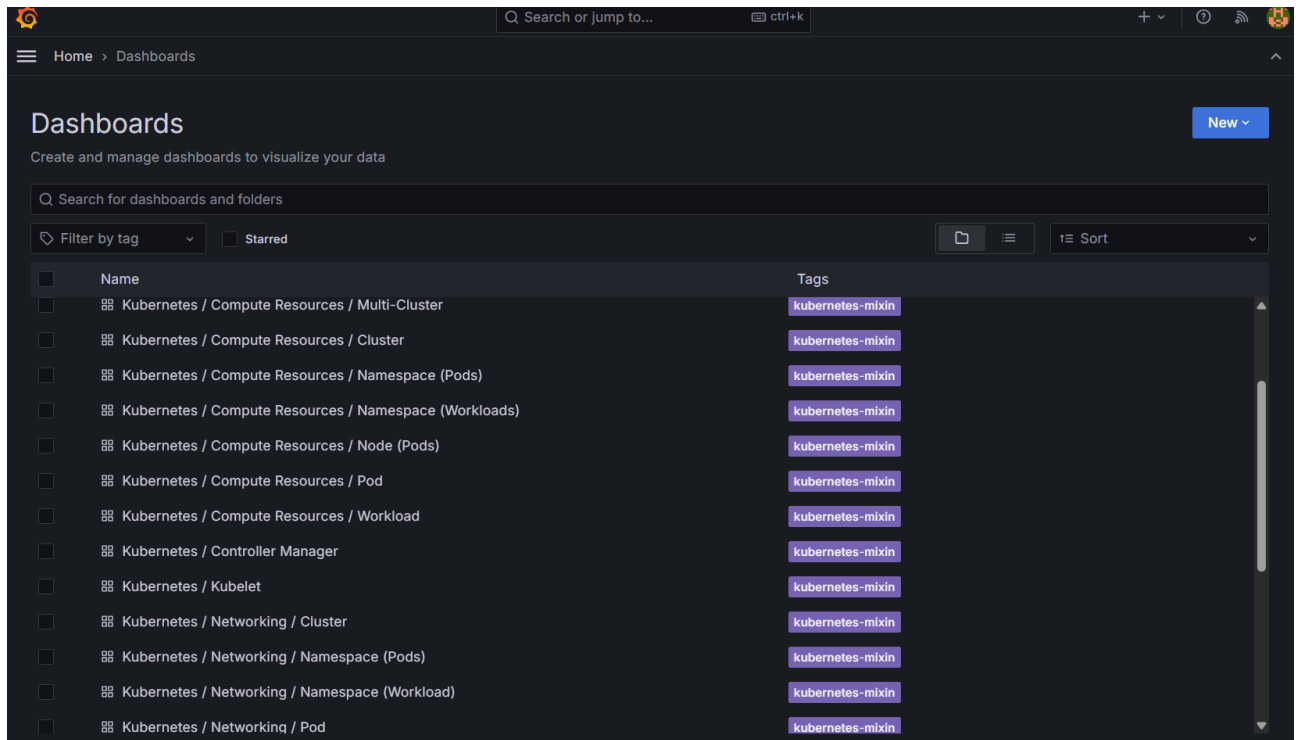


Figure 28 - Delete Integration/Sales

6. Observability

Deploy observability infrastructure within the Kubernetes cluster to monitor the health, performance, and availability of all services using Prometheus and Grafana



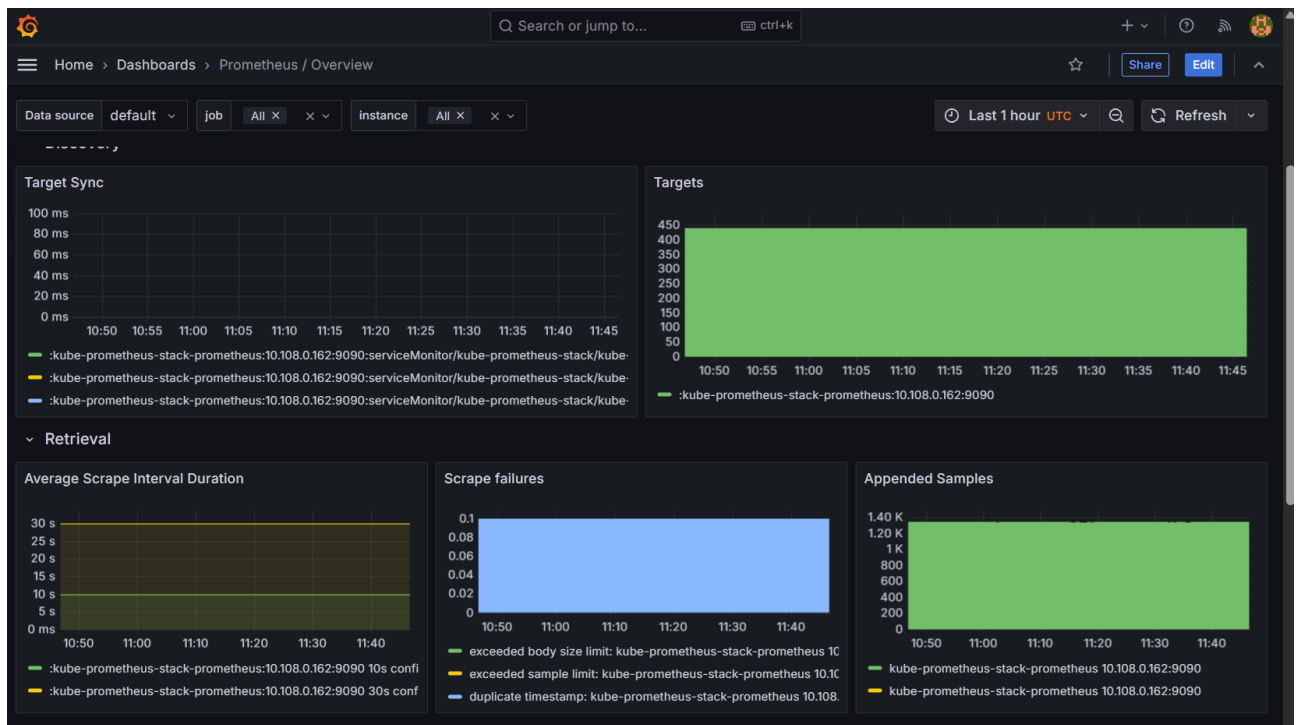
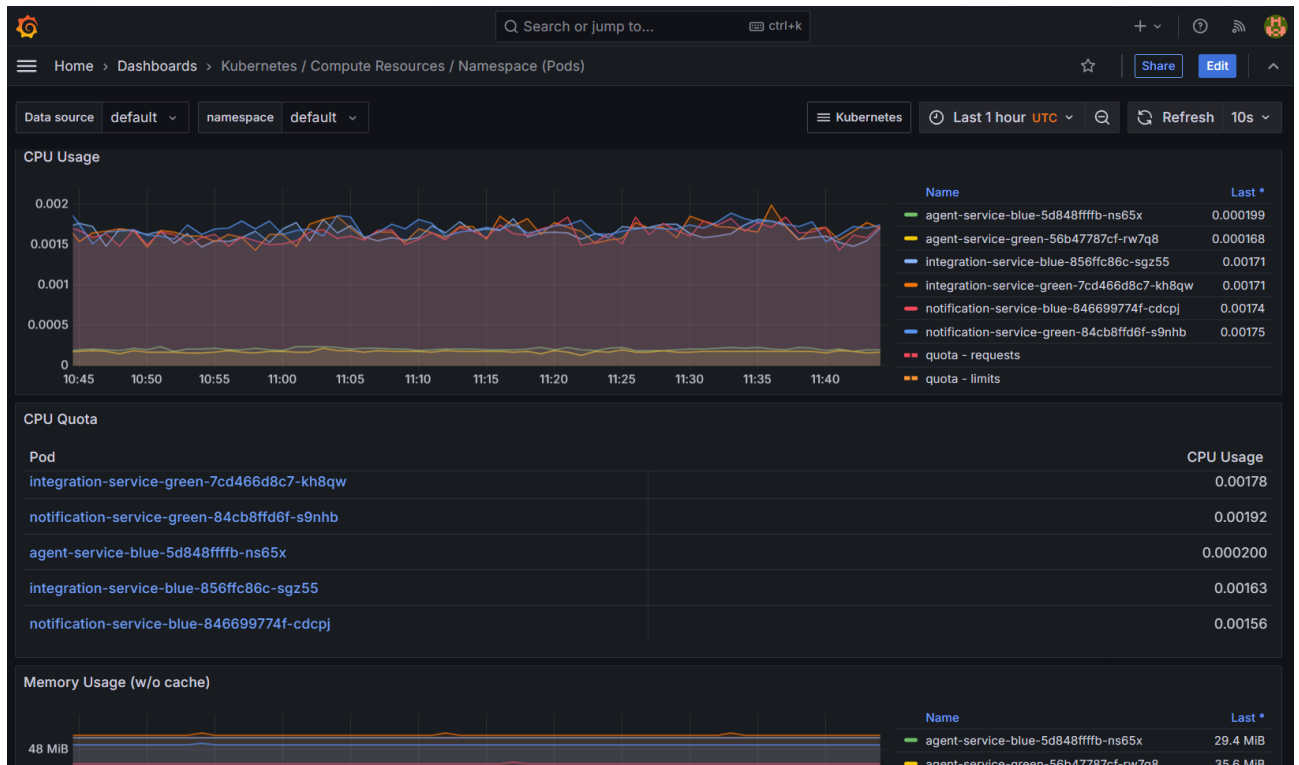


Figure 29 - Grafana-Dashboards

7. MoonInsurance System Deployment Runbook

7.1. Overview

The system comprises multiple microservices deployed in a DigitalOcean Kubernetes environment with CI/CD integration via GitHub Actions. It includes service containerization, automatic deployments, observability using Grafana, and visual analytics via AWS QuickSight.

Runbook name	MediTrack System Deployment Runbook
Runbook description	<p>Runbook to demonstrate the deployment and testing of the MoonInsurance system, which includes the following core components:</p> <ol style="list-style-type: none">1. Agent Microservice: Manages operations related to insurance agents, including adding, updating, retrieving, and removing agent records. This service plays a vital role in maintaining accurate agent profiles across branches.2. Integration Microservice: Handles all insurance product sales transactions, associating agents with products sold, sale dates, and revenue. It ensures accurate tracking of performance and branch-wise sales activity.3. Notification Microservice: Responsible for generating and delivering timely sales performance notifications to agents, helping stakeholders stay informed of milestones, targets, and achievements.4. Aggregator Microservice (Scheduled Job): Periodically aggregates data across services to calculate team performance, sales against targets, and branch-level metrics. It uploads this data to AWS Redshift, enabling visual analytics via AWS QuickSight.
Owner	Fouzul Hassan
Version	v1.0
Version date	16-04-2025
On this page	<ul style="list-style-type: none">- Prerequisites- Deployment steps- Testing procedures- Troubleshooting

7.2. Support Contacts

Expertise Level	Team	Contact
Developer and Owner	Fouzul Hassan	fouzul.20233214@iit.ac.lk

7.3. Process

Step	Task	Command/Action
Setting Up the Environment	Verify Prerequisites - Ensure Kubernetes cluster, DigitalOcean Container Registry, and GitHub CI/CD are set up.	Confirm infrastructure and access configurations.
	Configure Access - Authenticate with DigitalOcean and configure kubectl.	<code>doctl auth init</code> <code>doctl kubernetes cluster kubeconfig save <cluster-name></code>
Trigger the CI/CD Pipeline	Push Code Changes - Push updates to the main branch of GitHub.	<code>git add .</code> <code>git commit -m "Update application code"</code> <code>git push origin main</code>
	Monitor Workflow - Verify CI/CD workflow trigger in GitHub.	Open GitHub repository and navigate to Actions tab.
Blue Deployment	Switch Traffic - Route traffic to Blue Environment via Blue Ingress Controller.	<code>kubectl describe ingress blue- ingress -n meditrack</code>
Build and Push Docker Images	Build Docker Images -CI/CD pipeline builds Docker images.	Check CI/CD logs in GitHub Actions for build status.
	Push to Registry - Docker images are pushed to DigitalOcean Container Registry.	Validate via DigitalOcean Container Registry dashboard.
Green Deployment	Deploy Green - Deploy updated images to Green Environment in Kubernetes.	<code>kubectl get deployments -n meditrack-green</code>
	Run Tests - Execute automated unit and integration tests in Green Environment.	<code>kubectl logs <pod-name> -n meditrack-green</code>
Switch to Green Environment	Deploy to Blue - Update Blue Environment with the stable deployment.	<code>kubectl get pods -n meditrack</code>
	Switch Traffic - Ingress Controller routes traffic to Green Environment.	Confirm ingress rules using <code>kubectl describe ingress</code> .
Post-Deployment Testing	Testing APIs - Use Postman to send API requests to the Green Environment.	<code><green-ingress- ip>:<port>/api/<endpoint></code>

	Validate Results - Check logs for errors and ensure proper responses.	<i>kubectrl logs <pod-name> -n meditrack-green</i>
Fallback Mechanism	Handle Failures - Halt deployment and notify developers if Green Environment fails.	CI/CD pipeline logs indicate failure and notification to developers.
	Rollback Changes - Roll back Blue Environment if required.	kubectrl rollout undo deployment <i><deployment-name> -n meditrack</i>
	Retry Deployment - Fix issues, push updates, and retrigger CI/CD.	Follow Step 2 to restart the process.

8. Dashboard – MoonInsurance System

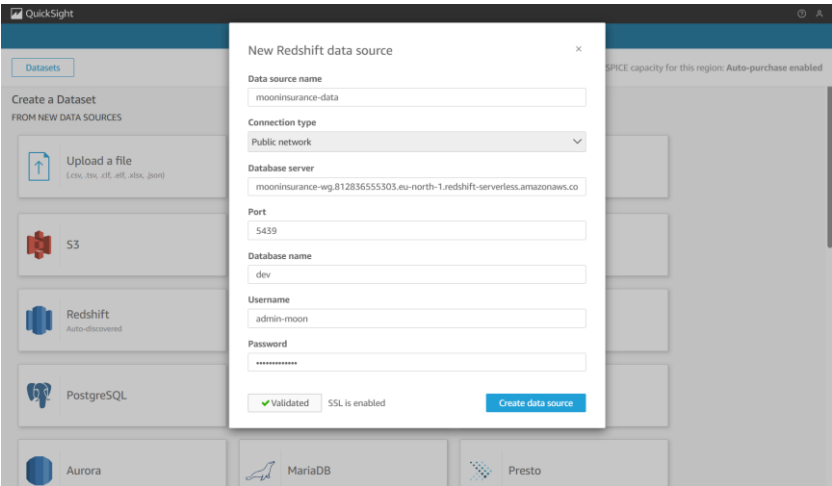


Figure 30 - Connecting-MoonInsurance

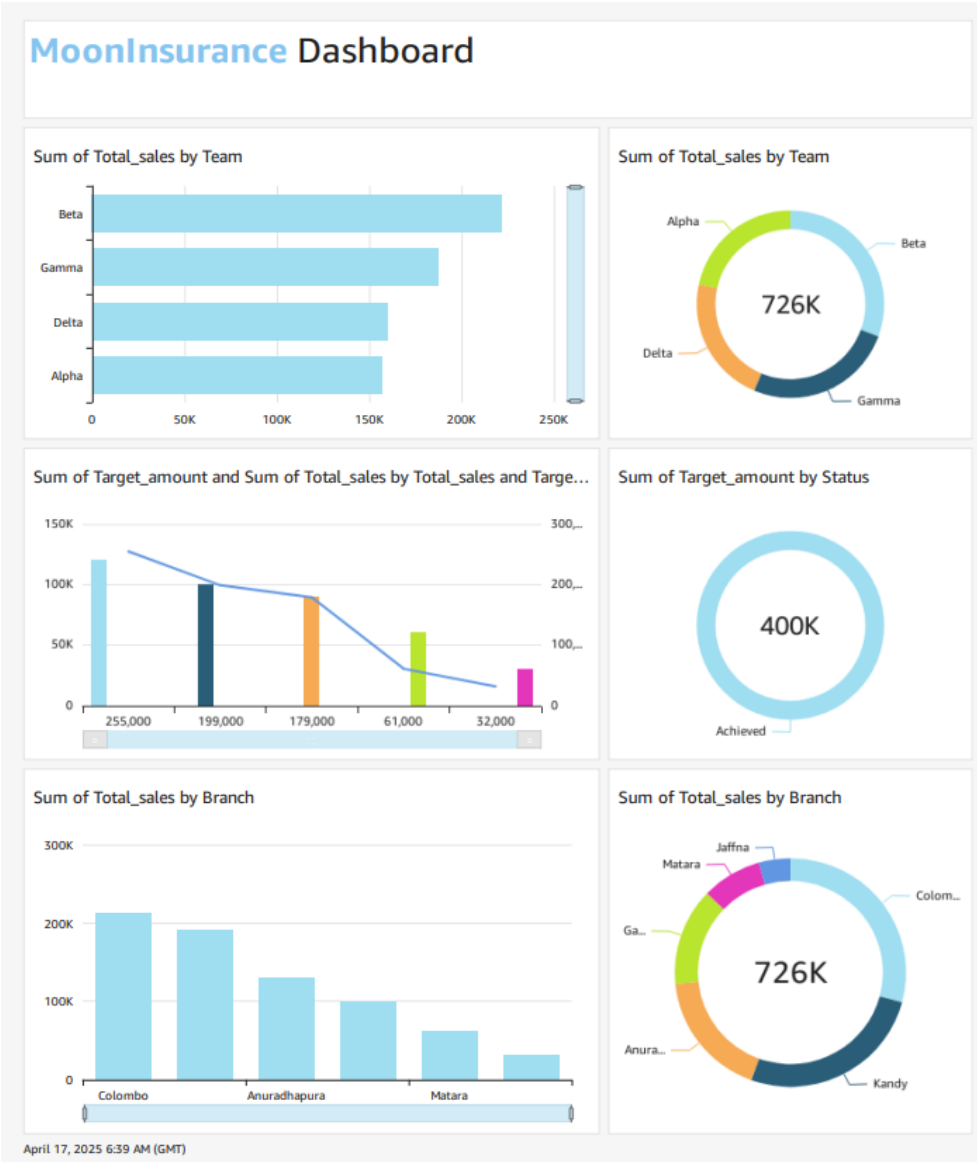


Figure 31 – Dashboard

9. Metrics and High-Performance Analytics Support

Three Key Metrics Considered:

1. Total Sales by Team
2. Total Sales by Branch
3. Target Amount vs. Total Sales by Product

Analytics & Visualization Benefits:

- These metrics are **stored in AWS Redshift**, enabling **scalable, high-performance analytics** across large volumes of sales data.
- Visualized via **AWS QuickSight**, these interactive dashboards allow administrators to:
 - **Identify top-performing teams and branches**
 - **Monitor progress against product targets**
 - **Gain actionable insights** to adjust sales strategies and resource allocation
- The **real-time dashboards** enhance operational decision-making and support strategic planning for insurance corporate providers.